



Parallelization in R

Irina Degtiar¹ & Anna Zink²

¹ Department of Biostatistics, Harvard T. H. Chan School of Public Health

² Department of Health Care Policy, Harvard Medical School



Help! My code's too slow :(

- How come?
 - Profiling & benchmarking
- How do I fix it?
 - Strategies
- How do I parallelize in R?
 - Hardware & terminology
 - Data vs. task parallelization
 - Forking vs. sockets
 - Packages in R
 - Example

Acknowledgements

- Christine Choirat: BST 262

Help! My code's too slow :(

- How come?
 - Profiling & benchmarking
- How do I fix it?
 - Strategies
- How do I parallelize in R?
 - Hardware & terminology
 - Data vs. task parallelization
 - Forking vs. sockets
 - Packages in R
 - Example

Profiling & Benchmarking

- **Profiling:** finding bottlenecks
 - Time
 - Memory
- **Benchmarking:** removing bottlenecks
 - Can optimize for total time, CPU time, memory usage, I/O, etc.
 - Tradeoff between runtime and researcher time
 - Run code multiple times to compare distributions

Profiling & Benchmarking: R

- **Profiling**

- Base R: `system.time()`, `Rprof()`, `Rprofmem()`
- R packages: `profvis`, `profr`, `lineprof`, `proftools`, `GUIProfiler`, `aprof`

- **Benchmarking**

- R packages: `microbenchmark`, `rbenchmark`

Help! My code's too slow :(

- How come?
 - Profiling & benchmarking
- How do I fix it?
 - Strategies
- How do I parallelize in R?
 - Hardware & terminology
 - Data vs. task parallelization
 - Forking vs. sockets
 - Packages in R
 - Example

Strategies to optimize memory

- More efficient code
 - Vectorize
 - Avoid series of scalars
 - E.g., `runif(n)` rather than `runif(1)` within loop
 - Avoid concatenating within loop
 - E.g., `x[i] <- new_value` rather than `x <- c(x, new_value)`
 - Note: subscripting often will create a copy but will be faster than growing the object
 - Avoid copies
 - R creates (temp) copies when mutating/modifying object
 - Only use what you need: `subset`, `rm()` objects you don't need, use iterators
- Better data storage
 - Matrices rather than `data.frame`
 - `Matrix`: sparse matrices
 - `data.table`
 - `bigmemory`, `ffbase`, `diskframe`, `bigstatsr`

Strategies to optimize time

- More efficient code
- Faster languages/packages
- Parallelization

Strategies to optimize time

- More efficient code
 - Vectorize
 - Use vectors not scalars
 - Avoid for loops, use `apply` family or `purrr::map()`
 - Use vectorized functions (e.g., `rowSums()`, `colMeans()`, matrix algebra)
 - Vectorize your own functions: `Vectorize`
 - Be specific
 - Use functions specific to your data type
 - E.g., `vapply` instead of `sapply`
 - Specify arguments to function
 - E.g., specify column type for `read.csv()` or levels for `factor()`
 - Pre-compile your compute functions: `compiler::cmpfun()`
- Faster languages/packages
- Parallelization

Strategies to optimize time

- More efficient code
- Faster languages/packages
 - Faster all-around: Rcpp
 - Faster I/O
 - `fst`: data.tables/data.frames
 - `feather`: data.tables/data.frames (Python, Spark compatibility)
 - `qs`: all objects
 - Trim down models: `strip`
 - Look for other existing packages, especially for your specific task
- Parallelization

Strategies to optimize time


- More efficient code
- Faster languages/packages
- Parallelization
 - Words of warning
 - Remember overhead: I/O + thread creation
 - Servers often have worse machines than your laptop
 - Can make memory issues worse

Help! My code's too slow :(

- How come?
 - Profiling & benchmarking
- How do I fix it?
 - Strategies
- How do I parallelize in R?
 - Hardware & terminology
 - Data vs. task parallelization
 - Forking vs. sockets
 - Packages in R
 - Example

Hardware & terminology

- Machine, node, cluster, server
 - Machine/node = single physical computer
 - Cluster = interconnected computers
 - Server: lets computers work together as cluster
- Core, processor, CPU, socket
 - Core = basic computational unit
 - Processor/CPU/socket: can have multiple cores
- Parallel vs. distributed computing
 - Same machine = parallel computing
 - Different machines (cluster) = distributed computing
- Master and workers
 - Master/parent: orchestrates processes across workers
 - Worker/child/slave: each of the cores, does the work



Data parallelization

- Divide data into partitions and perform same task on every partition
- Use cases
 - Data can't fit in memory
 - Data lives in different places (e.g., Sentinel)
 - Speed up computation when results combinable (e.g., bootstrap, sum or means)
- R
 - tidyverse split-apply-combine: `data %>% group_by() %>% summarize()`
 - MapReduce
 - Divide and conquer

Task parallelization

- Divide independent tasks across processors and perform different task on every data copy
- Use cases
 - Ensemble of algorithms (SuperLearner , random forests)
 - Comparison of statistical models

Forking vs. sockets in R

- Forking
 - Every worker inherits properties of existing R session (data, packages, etc.)
 - Data isn't copied until it's modified
 - Usually faster than sockets
 - Doesn't work on Windows (only Linux/Unix/OSX)
- Sockets
 - Every worker gets new version of R
 - Need to explicitly set up cluster, including loading packages: `cl <- makeCluster(4)`
 - Workers can communicate via network sockets
 - Works on Windows too

Popular Packages in R for Parallelization

- parallel
 - mclapply vs parLapply
- foreach
- doParallel

Beyond your laptop

- Software for parallelization
 - H2O
 - Spark
- Parallelization in the cloud
 - AWS, Google Compute Engine, Azure
- Parallelization on clusters
 - SLRM

Practicing Parallel Programming

Example code to get you started!

<https://github.com/zinka88/hpds-parallel>

Further resources

- Profiling
 - <http://adv-r.had.co.nz/Profiling.html>
- Fun read of R frustrations & solutions
 - http://www.burns-stat.com/pages/Tutor/R_inferno.pdf
- Parallelization
 - https://cran.r-project.org/web/views/HighPerformanceComputing.html?fbclid=IwAR3_RJ_09QcZTUz0MvCXRYOc3YZ_oTj1xcV9p_Gh4U0IyjId4u4QJeymfhE
 - <http://dept.stat.lsa.umich.edu/~jerrick/courses/stat701/notes/parallel.html>
 - <https://www.r-bloggers.com/2018/09/simple-parallel-processing-in-r/>
 - <https://bookdown.org/rdpeng/rprogdatascience/parallel-computation.html>